

x86 Assembly Language

Shinwoo Kim

Teaching Assistant

shinwookim@pitt.edu

<https://www.pitt.edu/~shk148/>

Spring 2023, Term 2234

Friday 12 PM Recitation

5502 Sennott Square

Mar 2nd, 2023

Malloc Project Due Tonight!

Don't forget to:

- *Submit to Gradescope*
- *See Gradescope for late submissions deadline*
- *Schedule checkoffs* *after the break*

Course News

➤ Project 2 due tonight

- 11:59 PM
- Don't forget to submit to Gradescope!
- *See Gradescope for late submission date*

➤ Spring Break

- No class or recitation next week (duh)

➤ New lab

➤ Project 3 released!

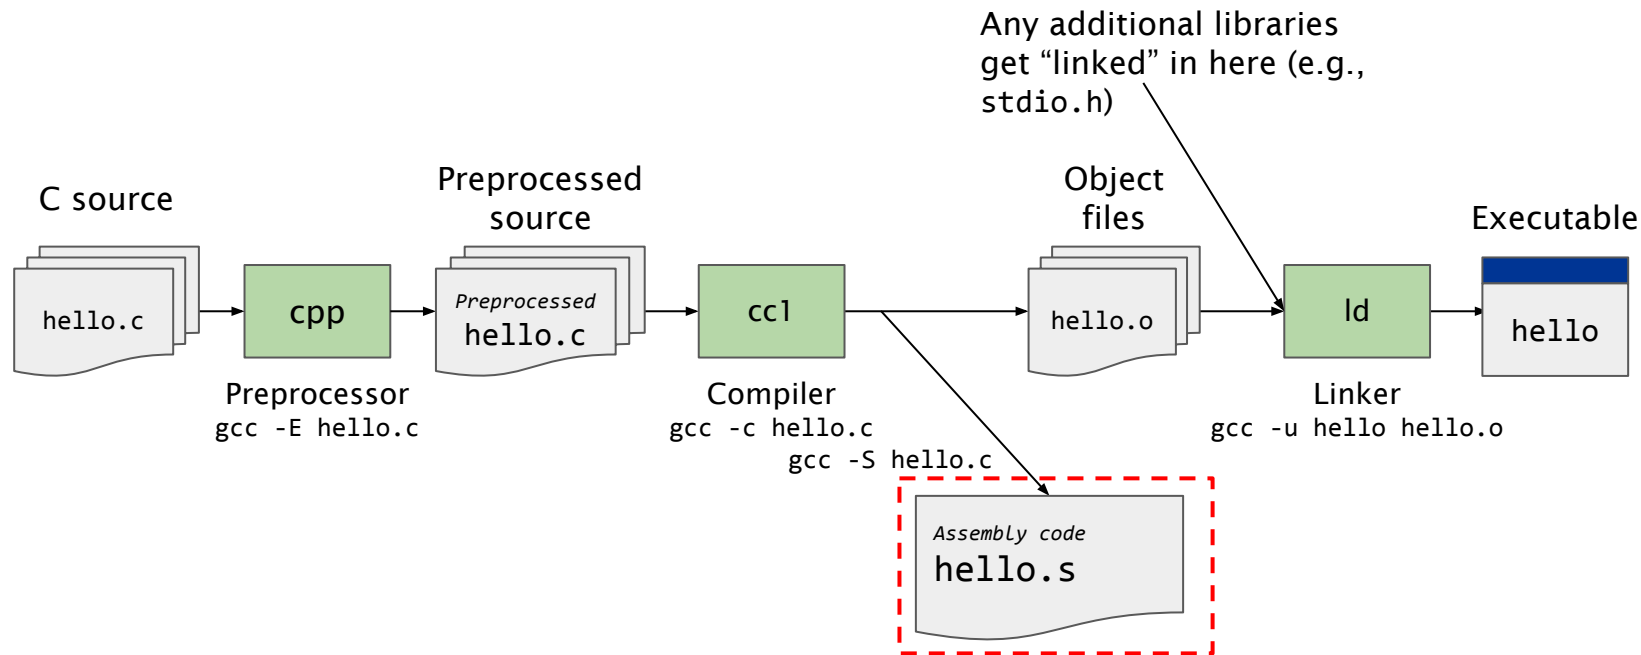
- Assembly Project
- 3 weeks to complete
 - *YOU ARE NOT EXPECTED TO WORK ON THE PROJECT DURING BREAK*

Assembly Language

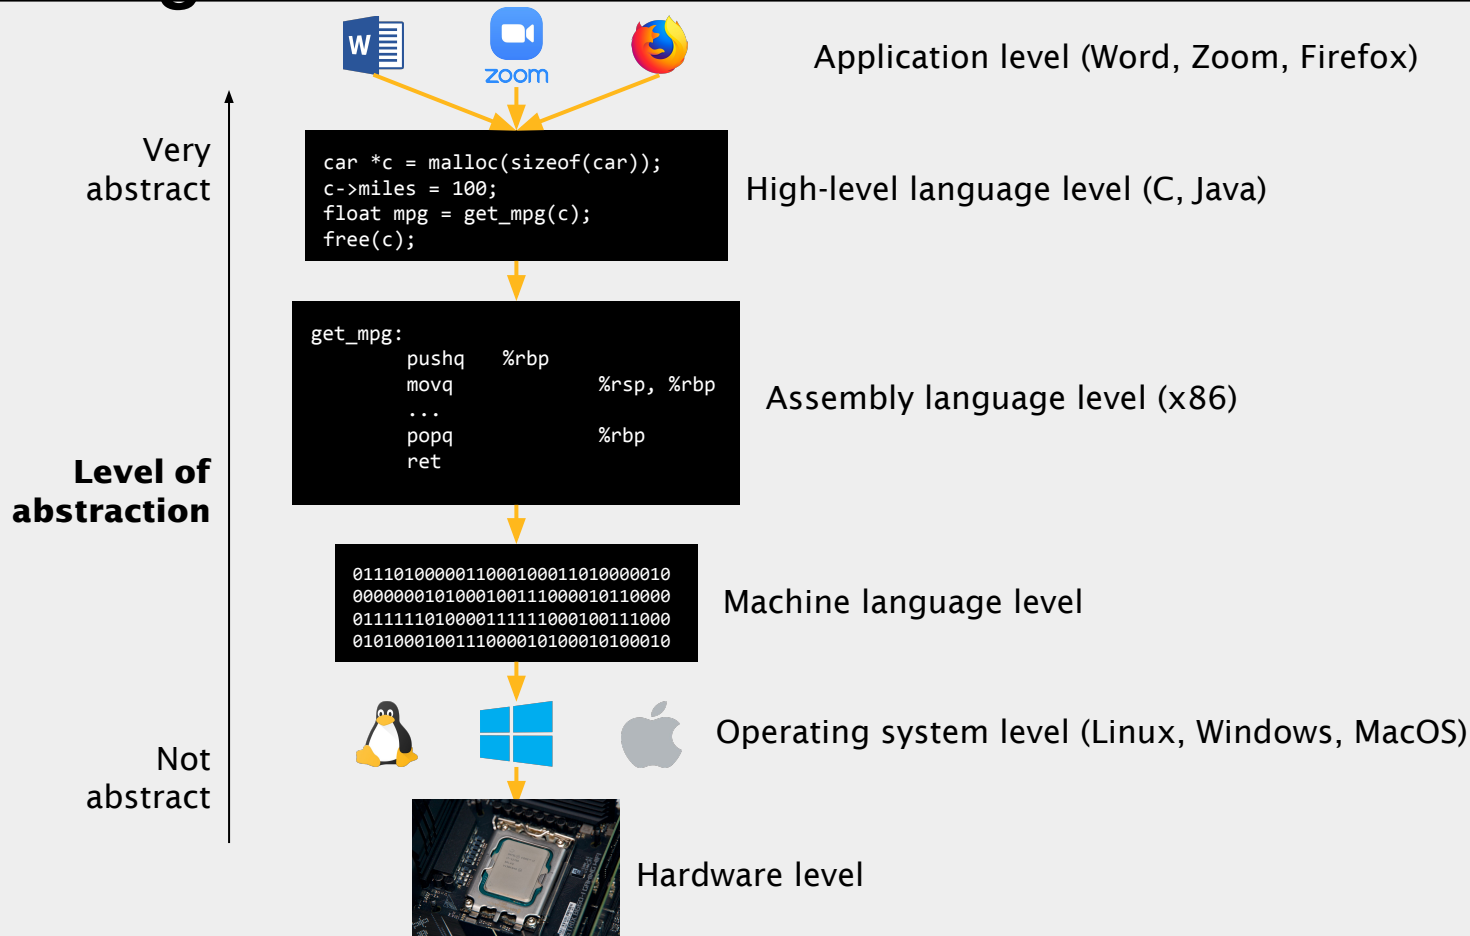
Because decoding 1s and 0s is hard

How a program is built

gcc hello.c



Moving down the ladder of abstractions



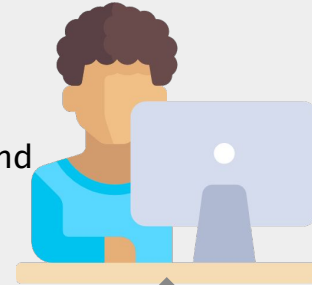
What is assembly?

→ **Assembly language** is a human-readable textual representation of machine language

High-level language
(C, Java)

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Relatively Easy for us to understand



```
get_mpg:  
    pushq    %rbp  
    movq     %rsp, %rbp  
    ...  
    popq     %rbp  
    ret
```

Assembly acts as a
translator between
high-level code and
machine code

Machine language

```
011101000001100010001101000  
001000000001010001001110000  
101100000111111010000111111  
100000111111010000111111100
```

Easy for computer
to understand

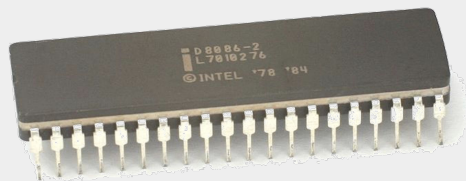


Enter x86

→ In CS447 Computer Organization & Assembly, you used **MIPS**

- ◆ Which was based on a Reduced Instruction Set Computer (**RISC**) ISA
 - Small number of instructions
 - Simple instructions

→ Now, we will use the **x86 asm (CISC)**



Intel 8086
Released 1978



Intel i9-10900K
Released 2020

x86 assembly language

- Epitome of Complex Instruction Set Computer (**CISC**)
 - Lots of instructions and ways to use them
 - Hundreds of instructions
- Designed for humans to write
 - From way back when programmers used to program in assembly language
 - A time before compilers or high-level languages
- Complex (multi-step) instructions
 - Instruction to search a string for a character
 - **F2XM1** computes $2^x - 1$
 - Computes the exponential value of 2 to the power of the source operand minus 1. The source operand is located in register ST(0) and the result is also stored in ST(0). The value of the source operand must lie in the range - 1.0 to +1.0. If the source value is outside this range, the result is undefined.
- Fewer instructions to write the same program
 - compared to RISC

But why use asm, if I can just code in C?

- Any C source can be compiled to assembly
 - `gcc -S <SOURCE>.c`
 - Not *really* helpful
- But what if we don't have the source code?
 - such as a `.exe` program you downloaded from the web
- You can **disassemble** any compiled program to emit the assembly
- What can you do with this?
 - Examine behavior of a program
 - Reverse engineering!

But why use asm, if I can just code in C?

Assembly is **good** for:

- Understanding the machine
 - ◆ You get to see what exactly the CPU is doing
- Better optimization of routines
 - ◆ Think you're better than a compiler?
- Programming hardware-dependent routines
 - ◆ E.g., compilers, operating systems,...
- Reverse-engineering and code obfuscation
 - ◆ malware/driver analysis...

Knowing assembly will enhance your code!

Assembly is **bad** for:

- Portability is lost
 - ◆ Code only works for a particular architecture, or processor
- Obfuscate the code
 - ◆ Not everyone can read assembly
 - **But you can!**
- Debugging is hard
 - ◆ Most debuggers are lost when hitting assembly
 - **But not GDB!**
- Optimizations is tedious
 - ◆ Tbh, you can't beat a modern compiler

Use it with caution and sparsity!

One machine code, two assembly

- Assembly language is simply a textual representation of machine language
⇒ Multiple representations for the same machine language

AT&T Syntax

- Developed by AT&T (duh)
- Used by GNU Assembler (**gas**)
- Opcode appended by type:
 - b** – byte (8 bit)
 - w** – word (16 bit)
 - l** – long (32 bit)
 - q** – quad (64 bit)
- First operand is **source**
- Second operand is **destination**
- Dereferences are denoted by **()**

Intel Syntax

- Developed by Intel (duh)
- Used by Microsoft (MASM), intel, NASM
- Type sizes are spelled out:
 - BYTE** – 1 byte
 - WORD** – 2 bytes
 - DWORD** – 4 bytes (double word)
 - QWORD** – 8 bytes (quad word)
- First operand is **destination**
- Second operand is **source**
- Dereferences are denoted by **[]**

Keeping track of the registers

- Like in MIPS, x86 has calling conventions
 - The **C Application Binary Interface (ABI)**
 - Like MIPS, certain registers are typically used for returns values, args, etc
- The ABI is not defined by the language, but rather the OS
 - Windows and Linux (UNIX/System V) have a different C ABI
- In our x86-64 Linux C ABI,
 - `%rdi`, `%rsi`, `%rdx`, `%rcx`, `%r8`, `%r9` are used to pass arguments (like the `a` registers in MIPS)
 - Remaining arguments go on the stack
 - A function callee must preserve `%rbp`, `%rbx`, `%r12`, `%r13`, `%r14`, `%r15` (like the `s` registers in MIPS)
 - `%rax` (overflows into `%rdx` for 128-bits) stores the return value (like `v0`, `v1` in MIPS)
- Reference manual provides extra information

Will I have to write assembly code for this course?

- **No!** No matter how good you are at programming, you are no match for a modern compiler
 - Modern Compilers are just too good at optimization
 - There was a time when humans outperformed compilers
 - Those days are long gone now...
- However, you should be able to ***read*** assembly code
 - To figure out what your machine is doing
 - To *guess* the C code
- By the end of this lab, you should be able to freely translate assembly and C

Diving into the Code!

See code: <https://github.com/shinwookim/asm-demo>

Hello World! x86 edition

```
#include <stdio.h>
int main(void)
{
    puts("Hello World!");
    return 0;
}
```

text (code) segment:

```
55 48 89 E5 BF 00 00 00 00 E8 00 00 00
00 B8 00 00 00 00 5D C3
```

data segment:

```
48 65 6C 6C 6F 2C 20 57 6F 72 6C
```

```
// Symbol table and other info omitted
```

```
.LC0:
    .string "Hello World!"
main:
    pushq    %rbp
    movq     %rsp, %rbp # rsp = stack pointer
    movl     $.LC0, %edi # push func args
    call     puts # call a function
    movl     $0, %eax # eax = return register
    popq     %rbp # prepare to return
    ret      # return
```

Linker

Executable

Debugging Assembly

- Recall that **GDB** worked on *executables*
 - You ran `gdb mdriver` and not ~~`gdb mdriver.e`~~
- Having the source was nice
 - We used the `-g` flag when compiling
 - which allowed us to use `layout src` to view the code during execution
- ...but not necessary
- What if we don't have a source file ? (or the program was compiled without `-g` flag)
 - We can still run GDB!
 - Won't be able to see the source code ⇒ We need to inspect assembly code

Reading symbols from a.out...

(No debugging symbols found in a.out)

Displaying the assembly with **disas**

- Suppose we are in paused in a breakpoint
- We can view the assembly code around our current memory address using **disas**
 - Memory address that is held by the program counter
- But how do we set a breakpoint
 - if we don't have the code?
- Surely, we need a way to view ASM
 - Without first setting a breakpoint right?

```
Dump of assembler code for function __GI_IO_puts:
Address range 0x7ffff7e09ed0 to 0x7ffff7e0a069:
=> 0x00007ffff7e09ed0 <+0>:      endbr64
0x00007ffff7e09ed4 <+4>:      push    %r14
0x00007ffff7e09ed6 <+6>:      push    %r13
0x00007ffff7e09ed8 <+8>:      push    %r12
0x00007ffff7e09eda <+10>:     mov     %rdi,%r12
0x00007ffff7e09edd <+13>:     push    %rbp
0x00007ffff7e09ede <+14>:     push    %rbx
0x00007ffff7e09edf <+15>:     sub     $0x10,%rsp
0x00007ffff7e09ee3 <+19>:     call   0x7ffff7db1490 <__ABS*+0xa8720@plt>
0x00007ffff7e09ee8 <+24>:     mov     0x197f49(%rip),%r13      # 0x7ffff7fa1e38
0x00007ffff7e09eef <+31>:     mov     %rax,%rbx
0x00007ffff7e09ef2 <+34>:     mov     0x0(%r13),%rbp
0x00007ffff7e09ef6 <+38>:     mov     0x0(%rbp),%eax
0x00007ffff7e09ef9 <+41>:     and     $0x8000,%eax
0x00007ffff7e09efe <+46>:     jne     0x7ffff7e09f58 <__GI_IO_puts+136>
0x00007ffff7e09f00 <+48>:     mov     %fs:0x10,%r14
0x00007ffff7e09f09 <+57>:     mov     0x88(%rbp),%r8
0x00007ffff7e09f10 <+64>:     cmp     %r14,0x8(%r8)
0x00007ffff7e09f14 <+68>:     je      0x7ffff7e0a008 <__GI_IO_puts+312>
0x00007ffff7e09f1a <+74>:     mov     $0x1,%edx
0x00007ffff7e09f1f <+79>:     lock cmpxchg %edx,(%r8)
0x00007ffff7e09f24 <+84>:     jne     0x7ffff7e0a050 <__GI_IO_puts+384>
0x00007ffff7e09f2a <+90>:     mov     0x88(%rbp),%r8
0x00007ffff7e09f31 <+97>:     mov     0x0(%r13),%rdi
0x00007ffff7e09f35 <+101>:    mov     %r14,0x8(%r8)
0x00007ffff7e09f39 <+105>:    mov     0xc0(%rdi),%eax
--Type <RET> for more, q to quit, c to continue without paging--
```

Displaying the assembly with `layout asm`

- The `layout asm` command displays the assembly of the entire program
 - You can scroll through the code and identify the memory addresses to set breakpoints
- But what if your program is *Huuge?*
 - That's gonna be a lot of scrolling

```
0x1119 <__do_global_dtors_aux+25>    je      0x1127 <__do_global_dtors_aux+39>
0x111b <__do_global_dtors_aux+27>    mov     0x2ee6(%rip),%rdi             # 0x4008
0x1122 <__do_global_dtors_aux+34>    call   0x1040 <__cxa_finalize@plt>
0x1127 <__do_global_dtors_aux+39>    call   0x1090 <deregister_tm_clones>
0x112c <__do_global_dtors_aux+44>    movb    $0x1,0x2edd(%rip)           # 0x4010 <completed.0>
0x1133 <__do_global_dtors_aux+51>    pop     %rbp
0x1134 <__do_global_dtors_aux+52>    ret
0x1135 <__do_global_dtors_aux+53>    nopl    (%rax)
0x1138 <__do_global_dtors_aux+56>    ret
0x1139 <__do_global_dtors_aux+57>    nopl    0x0(%rax)
0x1140 <frame_dummy>                endbr64
0x1144 <frame_dummy+4>              jmp     0x10c0 <register_tm_clones>
0x1149 <main>                        endbr64
0x114d <main+4>                      push    %rbp
0x114e <main+5>                      mov     %rsp,%rbp
0x1151 <main+8>                      lea     0xeac(%rip),%rax             # 0x2004
```

exec No process In: L?? PC: ??
(gdb) |

Let's put the asm in a file ⇒ Now we can **ctrl+f**

objdump -d program > program.s

- GNU provides a tool called object dump for unix-like systems
 - Let's you inspect information from object files
 - The **-d** flag disassembles the program and displays the **.code** section
 - The **>** flag redirects your standard I/O output to a file

```
USER@thoth:$ objdump -d a.out
a.out:      file format elf64-x86-64
Disassembly of section .init:
0000000000001000 <_init>:
   1000:  f3 0f 1e fa                endbr64
   1004:  48 83 ec 08                sub     $0x8,%rsp
   1008:  48 8b 05 d9 2f 00 00      mov     0x2fd9(%rip),%rax        # 3fe8
   100f:  48 85 c0                  test    %rax,%rax
   1012:  74 02                     je      1016 <_init+0x16>
   1014:  ff d0                     call    *%rax
   1016:  48 83 c4 08                add     $0x8,%rsp
   101a:  c3                       ret
...
```

GDB Assembly Edition

- Back to GDB...
- You can still set **breakpoints**
 - Not at specific lines of code...but at specific instructions (which are stored in memory)
 - `break *0x00005555555515b`
 - Why the `*`?
 - `*main+24`
 - You can set breakpoints at function offsets
 - Get this from GDB's `layout asm`
- You can still step through your code
 - Again, not stepping through lines of code, but through CPU instructions
 - Using `stepi` instead of `step`
 - `nexti` instead of `next`
 - `Continue`

GDB Assembly Edition

- Examining Memory

- We can print values stored at memory address or at registers
- `print/format expr`
 - Indicate registers with `$` (NOT `%`)
 - To print a value stored in a memory address use `*`
 - `format` tells us how to interpret values at that memory location
 - `d`: decimal
 - `x`: hex
 - `t`: binary
 - `f`: floating point
 - `i`: instruction
 - `c`: character
 - `p/x $rdi` displays the content at `%rdi` in a decimal format
 - Just because you print it as decimal does not mean that the value is a decimal
 - Interpretation of values depends on the context (which you need to provide)
- `info registers` lets you see all registers at once

Need help with GDB?

See (fmr) TA Gavin's GDB videos on Canvas!

C Control Structures → Assembly

```
#include <stdio.h>
```

```
int main(void)
```

```
{
```

```
    for (int i = 0; i < 10; i++)
```

```
    {
```

```
        printf("%d", i);
```

```
    }
```

```
    return 0;
```

```
}
```

0x0000000000001155 <+12>:	movl	\$0x0, -0x4(%rbp)
0x000000000000115c <+19>:	jmp	0x117b <main+50>
0x000000000000115e <+21>:	mov	-0x4(%rbp), %eax
0x0000000000001161 <+24>:	mov	%eax, %esi
0x0000000000001163 <+26>:	lea	0xe9a(%rip), %rax
0x000000000000116a <+33>:	mov	%rax, %rdi
0x000000000000116d <+36>:	mov	\$0x0, %eax
0x0000000000001172 <+41>:	call	0x1050 <printf@plt>
0x0000000000001177 <+46>:	addl	\$0x1, -0x4(%rbp)
0x000000000000117b <+50>:	cmpl	\$0x9, -0x4(%rbp)
0x000000000000117f <+54>:	jle	0x115e <main+21>

C Control Structures → Assembly

```
#include <stdio.h>
```

```
int main(void)
```

```
{
```

```
    int i = 0;
```

```
    while (i < 10)
```

```
    {
```

```
        printf("%d", i);
```

```
        i++;
```

```
    }
```

```
    return 0;
```

```
}
```

0x00000000000001155	<+12>:	movl	\$0x0, -0x4(%rbp)
0x0000000000000115c	<+19>:	jmp	0x117b <main+50>
0x0000000000000115e	<+21>:	mov	-0x4(%rbp), %eax
0x00000000000001161	<+24>:	mov	%eax, %esi
0x00000000000001163	<+26>:	lea	0xe9a(%rip), %rax
0x0000000000000116a	<+33>:	mov	%rax, %rdi
0x0000000000000116d	<+36>:	mov	\$0x0, %eax
0x00000000000001172	<+41>:	call	0x1050 <printf@plt>
0x00000000000001177	<+46>:	addl	\$0x1, -0x4(%rbp)
0x0000000000000117b	<+50>:	cmpl	\$0x9, -0x4(%rbp)
0x0000000000000117f	<+54>:	jle	0x115e <main+21>

C Control Structures → Assembly

```
#include <stdio.h>
```

```
int main(void)
```

```
{
```

```
    for (int i = 0; i < 10; i++)
```

```
    {
```

```
        printf("%d", i);
```

```
    }
```

```
    return 0;
```

```
}
```

```
0x00000000000001155 <+12>: movl    $0x0, -0x4(%rbp)
0x0000000000000115c <+19>: jmp     0x117b <main+50>
0x0000000000000115e <+21>: mov     -0x4(%rbp), %eax
0x00000000000001161 <+24>: mov     %eax, %esi
0x00000000000001163 <+26>: lea     0xe9a(%rip), %rax
0x0000000000000116a <+33>: mov     %rax, %rdi
0x0000000000000116d <+36>: mov     $0x0, %eax
0x00000000000001172 <+41>: call    0x1050 <printf@plt>
0x00000000000001177 <+46>: addl    $0x1, -0x4(%rbp)
0x0000000000000117b <+50>: cmpl    $0x9, -0x4(%rbp)
0x0000000000000117f <+54>: jle     0x115e <main+21>
```

Wait....why is the assembly code the same?

for loops == while loops!

Your CPU treats them the same way!

* do-while loops also work the same way (Write a short program and inspect the assembly!)

C Control Structures → Assembly

```
#include <stdio.h>
int main(void)
{
    int input;
    scanf("%d", &input);
    if (input > 10) printf("Big");
    else printf("Not Big");
    return 0;
}
```

11bf: 8b 45 f4	mov	-0xc(%rbp),%eax
11c2: 83 f8 0a	cmp	\$0xa,%eax
11c5: 7e 16	jle	11dd <main+0x54>
11c7: 48 8d 05 39 0e 00 00	lea	0xe39(%rip),%rax
11ce: 48 89 c7	mov	%rax,%rdi
11d1: b8 00 00 00 00	mov	\$0x0,%eax
11d6: e8 a5 fe ff ff	call	1080 <printf@plt>
11db: eb 14	jmp	11f1 <main+0x68>
11dd: 48 8d 05 27 0e 00 00	lea	0xe27(%rip),%rax
11e4: 48 89 c7	mov	%rax,%rdi
11e7: b8 00 00 00 00	mov	\$0x0,%eax
11ec: e8 8f fe ff ff	call	1080 <printf@plt>

Conditional statements works as expected

Who knew that `if-else` executed different based on
conditions?

Our *real* first assembly code analysis

Looking through a real program!

Special thanks to Jake Kasper for providing slides & code

C Control Structures → Assembly

```
#include <stdio.h>
```

```
int main(int argc, char **argv)
{
    int myNum = increment(5);
    printf("My num is %d\n", myNum);
    return 0;
}
```

```
int increment(int num)
{
    return ++num;
}
```

Prefix increment
Increments first, then returns

0000000000001149 <main>:

1149: f3 0f 1e fa

endbr64

114d: 55

push %rbp

114e: 48 89 e5

mov %rsp,%rbp

1151: 48 83 ec 20

sub \$0x20,%rsp

1155: 89 7d ec

mov %edi,-0x14(%rbp)

1158: 48 89 75 e0

mov %rsi,-0x20(%rbp)

115c: bf 05 00 00 00

mov \$0x5,%edi

1161: e8 23 00 00 00

call 1189<increment>

1166: 89 45 fc

mov %eax,-0x4(%rbp)

(...)

C Control Structures → Assembly

```
#include <stdio.h>
```

```
int main(int argc, char **argv)
```

```
{  
    int myNum = increment(5);  
    printf("My num is %d\n", myNum);  
    return 0;  
}
```

```
int increment(int num)  
{  
    return ++num;  
}
```

0000000000001189 <increment>:

1189: f3 0f 1e fa

endbr64

118d: 55

push %rbp

118e: 48 89 e5

mov %rsp,%rbp

1191: 89 7d fc

mov %edi,-0x4(%rbp)

1194: 83 45 fc 01

addl \$0x1,-0x4(%rbp)

1198: 8b 45 fc

mov -0x4(%rbp),%eax

119b: 5d

pop %rbp

119c: c3

ret

C Control Structures → Assembly

```
#include <stdio.h>
```

```
int main(int argc, char **argv)
```

```
{  
    int myNum = increment(5);  
    printf("My num is %d\n", myNum);  
    return 0;  
}
```

```
int increment(int num)
```

```
{  
    return ++num;  
}
```

`%rbp` needs maintains the current stack frame

- To preserve the previous stack frame
- it gets pushed onto the stack

00000000000001189 <increment>:

1189: f3 0f 1e fa

endbr64

118d: 55

push %rbp

118e: 48 89 e5

mov %rsp,%rbp

1191: 89 7d fc

mov %edi,-0x4(%rbp)

1194: 83 45 fc 01

addl \$0x1,-0x4(%rbp)

1198: 8b 45 fc

mov -0x4(%rbp),%eax

119b: 5d

pop %rbp

119c: c3

ret

C Control Structures → Assembly

```
#include <stdio.h>
```

```
int main(int argc, char **argv)
```

```
{  
    int myNum = increment(5);  
    printf("My num is %d\n", myNum);  
    return 0;  
}
```

```
int increment(int num)
```

```
{  
    return ++num;  
}
```

%edi is our first argument register, so we're moving the value of our argument (num) into the current stack frame
Why -0x4?

00000000000001189 <increment>:

1189:f3 0f 1e fa	endbr64
118d:55	push %rbp
118e:48 89 e5	mov %rsp,%rbp
1191:89 7d fc	mov %edi,-0x4(%rbp)
1194:83 45 fc 01	addl \$0x1,-0x4(%rbp)
1198:8b 45 fc	mov -0x4(%rbp),%eax
119b:5d	pop %rbp
119c:c3	ret

C Control Structures → Assembly

```
#include <stdio.h>
```

```
int main(int argc, char **argv)
{
    int myNum = increment(5);
    printf("My num is %d\n", myNum);
    return 0;
}
```

```
int increment(int num)
{
    return ++num;
}
```

Increment the value of the argument we just stored in the stack

00000000000001189 <increment>:

1189:f3 0f 1e fa	endbr64
118d:55	push %rbp
118e:48 89 e5	mov %rsp,%rbp
1191:89 7d fc	mov %edi,-0x4(%rbp)
1194:83 45 fc 01	addl \$0x1,-0x4(%rbp)
1198:8b 45 fc	mov -0x4(%rbp),%eax
119b:5d	pop %rbp
119c:c3	ret

C Control Structures → Assembly

```
#include <stdio.h>
```

```
int main(int argc, char **argv)
{
    int myNum = increment(5);
    printf("My num is %d\n", myNum);
    return 0;
}
```

```
int increment(int num)
{
    return ++num;
}
```

Move our data we've been editing in the stack, to our return register

```
00000000000001189 <increment>:
1189: f3 0f 1e fa                endbr64
118d: 55                          push %rbp
118e: 48 89 e5                    mov %rsp,%rbp
1191: 89 7d fc                    mov %edi,-0x4(%rbp)
1194: 83 45 fc 01                addl $0x1,-0x4(%rbp)
1198: 8b 45 fc                    mov -0x4(%rbp),%eax
119b: 5d                          pop %rbp
119c: c3                          ret
```

C Control Structures → Assembly

```
#include <stdio.h>
```

```
int main(int argc, char **argv)
{
    int myNum = increment(5);
    printf("My num is %d\n", myNum);
    return 0;
}
```

```
int increment(int num)
{
    return ++num;
}
```

Pop the stack frame from the stack, as we're about to return from the current function scope, and this will load the previous stack frame back to `%rbp`

00000000000001189 <increment>:

1189: f3 0f 1e fa	endbr64
118d: 55	push %rbp
118e: 48 89 e5	mov %rsp,%rbp
1191: 89 7d fc	mov %edi,-0x4(%rbp)
1194: 83 45 fc 01	addl \$0x1,-0x4(%rbp)
1198: 8b 45 fc	mov -0x4(%rbp),%eax
119b: 5d	pop %rbp
119c: c3	ret

C Control Structures → Assembly

```
#include <stdio.h>
```

```
int main(int argc, char **argv)
```

```
{  
    int myNum = increment(5);  
    printf("My num is %d\n", myNum);  
    return 0;  
}
```

```
int increment(int num)
```

```
{  
    return ++num;  
}
```

Return to caller

What about the return value?

It's already in the return register(%eax)

00000000000001189 <increment>:

1189: f3 0f 1e fa

endbr64

118d: 55

push %rbp

118e: 48 89 e5

mov %rsp,%rbp

1191: 89 7d fc

mov %edi,-0x4(%rbp)

1194: 83 45 fc 01

addl \$0x1,-0x4(%rbp)

1198: 8b 45 fc

mov -0x4(%rbp),%eax

119b: 5d

pop %rbp

119c: c3

ret

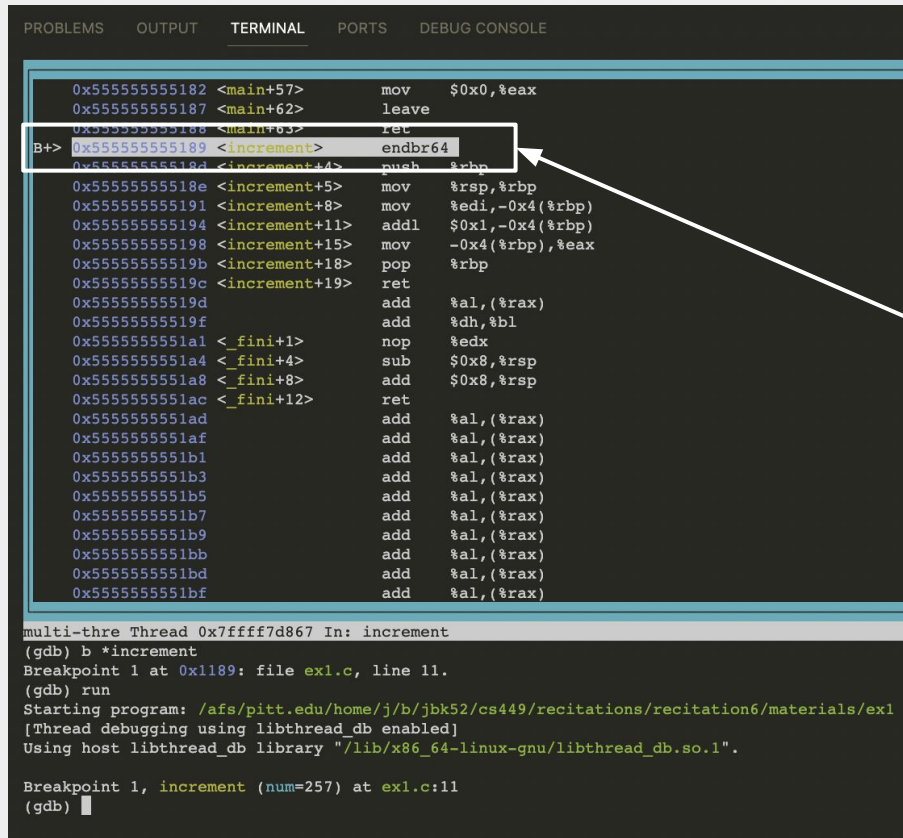
Let's inspect `increment()` with GDB

```
0x1149 <main>      endbr64
0x114d <main+4>     push  %rbp
0x114e <main+5>     mov   %rsp,%rbp
0x1151 <main+8>     sub   $0x20,%rsp
0x1155 <main+12>    mov   %edi,-0x14(%rbp)
0x1158 <main+15>    mov   %rsi,-0x20(%rbp)
0x115c <main+19>    mov   $0x5,%edi
0x1161 <main+24>    call  0x1189 <increment>
0x1166 <main+29>    mov   %eax,-0x4(%rbp)
0x1169 <main+32>    mov   -0x4(%rbp),%eax
0x116c <main+35>    mov   %eax,%esi
0x116e <main+37>    lea   0xe8f(%rip),%rax    # 0x2004
0x1175 <main+44>    mov   %rax,%rdi
0x1178 <main+47>    mov   $0x0,%eax
0x117d <main+52>    call  0x1050 <printf@plt>
0x1182 <main+57>    mov   $0x0,%eax
0x1187 <main+62>    leave
0x1188 <main+63>    ret
b+ 0x1189 <increment> endbr64
0x118d <increment+4> push  %rbp
0x118e <increment+5> mov   %rsp,%rbp
0x1191 <increment+8> mov   %edi,-0x4(%rbp)
0x1194 <increment+11> addl  $0x1,-0x4(%rbp)
0x1198 <increment+15> mov   -0x4(%rbp),%eax
0x119b <increment+18> pop   %rbp
0x119c <increment+19> ret

exec No process in:
(gdb) b *increment
Breakpoint 1 at 0x1189: file ex1.c, line 11.
(gdb)
```

Set a breakpoint at the start of the **assembly** for increment using the *****

Tracing through the code w/ GDB



The screenshot shows the GDB interface with the 'TERMINAL' tab selected. The assembly code is displayed, with a breakpoint set at address 0x55555555189, labeled '<increment>'. The instruction at this address is 'endbr64'. A white box highlights this instruction, and a white arrow points from a text box on the right to it. Below the assembly code, the GDB console shows the command 'multi-thre Thread 0x7ffff7d867 In: increment', followed by '(gdb) b *increment', 'Breakpoint 1 at 0x1189: file ex1.c, line 11.', '(gdb) run', and the program starting. The final line shows 'Breakpoint 1, increment (num=257) at ex1.c:11' and '(gdb) |'.

```
0x55555555182 <main+57>    mov     $0x0,%eax
0x55555555187 <main+62>    leave
0x55555555188 <main+63>    ret
B+> 0x55555555189 <increment>    endbr64
0x5555555518d <increment+4>    push    %rbp
0x5555555518e <increment+5>    mov     %rsp,%rbp
0x55555555191 <increment+8>    mov     %edi,-0x4(%rbp)
0x55555555194 <increment+11>   addl    $0x1,-0x4(%rbp)
0x55555555198 <increment+15>   mov     -0x4(%rbp),%eax
0x5555555519b <increment+18>   pop     %rbp
0x5555555519c <increment+19>   ret
0x5555555519d          add     %al,(%rax)
0x5555555519f          add     %dh,%bl
0x555555551a1 <_fini+1>         nop
0x555555551a4 <_fini+4>         sub     $0x8,%rsp
0x555555551a8 <_fini+8>         add     $0x8,%rsp
0x555555551ac <_fini+12>        ret
0x555555551ad          add     %al,(%rax)
0x555555551af          add     %al,(%rax)
0x555555551b1          add     %al,(%rax)
0x555555551b3          add     %al,(%rax)
0x555555551b5          add     %al,(%rax)
0x555555551b7          add     %al,(%rax)
0x555555551b9          add     %al,(%rax)
0x555555551bb          add     %al,(%rax)
0x555555551bd          add     %al,(%rax)
0x555555551bf          add     %al,(%rax)

multi-thre Thread 0x7ffff7d867 In: increment
(gdb) b *increment
Breakpoint 1 at 0x1189: file ex1.c, line 11.
(gdb) run
Starting program: /afs/pitt.edu/home/j/b/jbk52/cs449/recitations/recitation6/materials/ex1
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Breakpoint 1, increment (num=257) at ex1.c:11
(gdb) |
```

After running, we've hit the breakpoint at increment

Let's read the assembly line by line using **ni** (next instruction), though we can skip ahead a few lines until we get to the more important function details

Tracing through the code w/ GDB

```
0x55555555182 <main+57>    mov    $0x0,%eax
0x55555555187 <main+62>    leave
0x55555555188 <main+63>    ret
0x55555555189 <increment>    endbr64
0x5555555518d <increment+4>    push   %rbp
> 0x5555555518e <increment+5>    mov    %rsp,%rbp
0x55555555191 <increment+8>    mov    %edi,-0x4(%rbp)
0x55555555194 <increment+11>    addl   $0x1,-0x4(%rbp)
0x55555555198 <increment+15>    mov    -0x4(%rbp),%eax
0x5555555519b <increment+18>    pop    %rbp
0x5555555519c <increment+19>    ret
0x5555555519d      add    %al,(%rax)
0x5555555519f      add    %dh,%bl
0x555555551a1 <_fini+1>    nop    %edx
0x555555551a4 <_fini+4>    sub    $0x8,%rsp
0x555555551a8 <_fini+8>    add    $0x8,%rsp
0x555555551ac <_fini+12>   ret
0x555555551ad      add    %al,(%rax)
0x555555551af      add    %al,(%rax)
0x555555551b1      add    %al,(%rax)
0x555555551b3      add    %al,(%rax)
0x555555551b5      add    %al,(%rax)
0x555555551b7      add    %al,(%rax)
0x555555551b9      add    %al,(%rax)
0x555555551bb      add    %al,(%rax)
0x555555551bd      add    %al,(%rax)
0x555555551bf      add    %al,(%rax)
```

This is the line in which our stack frame pointer, `%rbp`, is being updated to contain the current stack address

Tracing through the code w/ GDB

```
0x55555555182 <main+57>    mov     $0x0,%eax
0x55555555187 <main+62>    leave
0x55555555188 <main+63>    ret
B+ 0x55555555189 <increment>    endbr64
0x5555555518d <increment+4>    push    %rbp
0x5555555518e <increment+5>    mov     %rsp,%rbp
> 0x55555555191 <increment+8>    mov     %edi,-0x4(%rbp)
0x55555555194 <increment+11>    addl    $0x1,-0x4(%rbp)
0x55555555198 <increment+15>    mov     -0x4(%rbp),%eax
0x5555555519b <increment+18>    pop     %rbp
0x5555555519c <increment+19>    ret
0x5555555519d <_fini+1>    add     %al,(%rax)
0x5555555519f <_fini+4>    add     %dh,%bl
0x555555551a1 <_fini+8>    nop     %edx
0x555555551a4 <_fini+12>   sub     $0x8,%rsp
0x555555551a8 <_fini+16>   add     $0x8,%rsp
0x555555551ac <_fini+20>   ret
0x555555551ad <_fini+24>   add     %al,(%rax)
0x555555551af <_fini+28>   add     %al,(%rax)
0x555555551b1 <_fini+32>   add     %al,(%rax)
0x555555551b3 <_fini+36>   add     %al,(%rax)
0x555555551b5 <_fini+40>   add     %al,(%rax)
0x555555551b7 <_fini+44>   add     %al,(%rax)
0x555555551b9 <_fini+48>   add     %al,(%rax)
0x555555551bb <_fini+52>   add     %al,(%rax)
0x555555551bd <_fini+56>   add     %al,(%rax)
0x555555551bf <_fini+60>   add     %al,(%rax)
```

We've now executed the instruction to add the current stack pointer to %rbp

We are also about to execute the line to put the argument register's contents into the stack frame, so let's check the value of the argument register:

p \$rdi →

```
(gdb) p $rdi
$1 = 5
```

This makes sense, as we passed 5 into our function in our C code

```
increment(5);
```

Tracing through the code w/ GDB

```
B+ 0x55555555189 <increment> endbr64
0x5555555518d <increment+4> push    %rbp
0x5555555518e <increment+5> mov     %rsp,%rbp
0x55555555191 <increment+8> mov     %edi,-0x4(%rbp)
> 0x55555555194 <increment+11> addl    $0x1,-0x4(%rbp)
0x55555555198 <increment+15> mov     -0x4(%rbp),%eax
0x5555555519b <increment+18> pop     %rbp
0x5555555519c <increment+19> ret
0x5555555519d add     %al,(%rax)
0x5555555519f add     %dh,%bl
0x555555551a1 <_fini+1> nop     %edx
0x555555551a4 <_fini+4> sub     $0x8,%rsp
0x555555551a8 <_fini+8> add     $0x8,%rsp
0x555555551ac <_fini+12> ret
0x555555551ad add     %al,(%rax)
0x555555551af add     %al,(%rax)
0x555555551b1 add     %al,(%rax)
0x555555551b3 add     %al,(%rax)
0x555555551b5 add     %al,(%rax)
0x555555551b7 add     %al,(%rax)
0x555555551b9 add     %al,(%rax)
0x555555551bb add     %al,(%rax)
0x555555551bd add     %al,(%rax)
0x555555551bf add     %al,(%rax)
0x555555551c1 add     %al,(%rax)
0x555555551c3 add     %al,(%rax)
0x555555551c5 add     %al,(%rax)
```

Now we stored the argument register value into our stack frame. To check that this update actually changed our stack frame, let's print the integer that lies below the stack pointer:

x/-4bx \$rbp → Read the previous 4 bytes

```
(gdb) x/-4bx $rbp
0x7fffffffef18c: 0x05 0x00 0x00 0x00
```

x/-1w \$rbp → Read the previous word (word is the size of an integer)

```
(gdb) x/-1w $rbp
0x7fffffffef18c: 5
```

We can see both of these led us to the value 5 being stored in the stack frame

Tracing through the code w/ GDB

```
0x55555555182 <main+57>    mov     $0x0,%eax
0x55555555187 <main+62>    leave
0x55555555188 <main+63>    ret
B+ 0x55555555189 <increment>  endbr64
0x5555555518d <increment+4>  push    %rbp
0x5555555518e <increment+5>  mov     %rsp,%rbp
0x55555555191 <increment+8>  mov     %edi,-0x4(%rbp)
0x55555555194 <increment+11> addl    $0x1,-0x4(%rbp)
> 0x55555555198 <increment+15> mov     -0x4(%rbp),%eax
0x5555555519b <increment+18> pop     %rbp
0x5555555519c <increment+19> ret
0x5555555519d      add     %al,(%rax)
0x5555555519f      add     %dh,%bl
0x555555551a1 <_fini+1>    nop     %edx
0x555555551a4 <_fini+4>    sub     $0x8,%rsp
0x555555551a8 <_fini+8>    add     $0x8,%rsp
0x555555551ac <_fini+12>   ret
0x555555551ad      add     %al,(%rax)
0x555555551af      add     %al,(%rax)
0x555555551b1      add     %al,(%rax)
0x555555551b3      add     %al,(%rax)
0x555555551b5      add     %al,(%rax)
0x555555551b7      add     %al,(%rax)
0x555555551b9      add     %al,(%rax)
0x555555551bb      add     %al,(%rax)
0x555555551bd      add     %al,(%rax)
0x555555551bf      add     %al,(%rax)
```

At this point, we've run the line to increment the value in the stack frame, and are waiting to execute this line.

To see if this change was made, let's again print out the values:

x/-4bx \$rbp → Read the previous 4 bytes as hex

```
(gdb) x/-4bx $rbp
0x7fffffffef18c: 0x06    0x00    0x00    0x00
```

x/-1wx \$rbp → Read the previous word (word is the size of an integer) as hex

```
(gdb) x/-1wx $rbp
0x7fffffffef18c: 0x00000006
```

Since the value changed to 6, the increment was successful, and we can see where that change occurred.

Tracing through the code w/ GDB

```
0x55555555182 <main+57>      mov     $0x0,%eax
0x55555555187 <main+62>      leave
0x55555555188 <main+63>      ret
B+ 0x55555555189 <increment>    endbr64
0x5555555518d <increment+4>    push    %rbp
0x5555555518e <increment+5>    mov     %rsp,%rbp
0x55555555191 <increment+8>    mov     %edi,-0x4(%rbp)
0x55555555194 <increment+11>   addl    $0x1,-0x4(%rbp)
0x55555555198 <increment+15>   mov     -0x4(%rbp),%eax
> 0x5555555519b <increment+18> pop     %rbp
0x5555555519c <increment+19> ret
0x5555555519d          add     %al,(%rax)
0x5555555519f          add     %dh,%bl
0x555555551a1 <_fini+1>          nop     %edx
0x555555551a4 <_fini+4>          sub     $0x8,%rsp
0x555555551a8 <_fini+8>          add     $0x8,%rsp
0x555555551ac <_fini+12>         ret
0x555555551ad          add     %al,(%rax)
0x555555551af          add     %al,(%rax)
0x555555551b1          add     %al,(%rax)
0x555555551b3          add     %al,(%rax)
0x555555551b5          add     %al,(%rax)
0x555555551b7          add     %al,(%rax)
0x555555551b9          add     %al,(%rax)
0x555555551bb          add     %al,(%rax)
0x555555551bd          add     %al,(%rax)
0x555555551bf          add     %al,(%rax)
```

%eax, the return register, should contain the value 6 that we want to return to the user. Let's see:

p \$rax → `(gdb) p $rax`
`$3 = 6`

%eax now contains the accurate return value from our function, so we can return to the previous caller after adjusting the stack.

Lab 4

Assembly Lab: ASM!

Now, it's your turn!

- In lab 4, you will practice:
 - Reading assembly
 - Recognizing common patterns
 - Using **gdb** to *debug assembly code + inspect memory!*
- Part A: Investigating the code!
 - Reading simple functions
 - Similar to what we just did
 - <https://godbolt.org/z/9c4Efqvoo>
 - Deep dive into *control flow, raise operations, hidden arguments*
 - **The Test.**
 - Can you read assembly code tell me what it does?
 - **Gradescope submission**
- Part B: Inspecting memory
 - Can you debug an executable by looking at assembly code and using gdb?
 - **Gradescope submission**

Works Referred

Jonathan Misurda's CS0449

Jake Kasper's CS 0449 Recitation Slides (Spring 2023)

Gavin Heinrichs-Majetich's CS 0449 Recitation Slides (Fall 2022)

Martha Dixon's CS 0449 Recitation Slides (Fall 2020)

Randal Bryant & David R. O'Hallaron's Computer Systems: A Programmer's Perspective

Carnegie Mellon University's 15-213: *Introduction to Computer Systems* (Fall 2017)