



Project 4:

Writing your own shell

Shinwoo Kim

Teaching Assistant

shinwookim@pitt.edu

<https://www.pitt.edu/~shk148/>

Spring 2023, Term 2234

Friday 12 PM Recitation

5502 Sennott Square

Mar 2nd, 2023

Course News!

▶ Exams

- Exam I grades were returned on March 24th, 2023
 - Check your email for class statistics
 - Request regrades if needed (this may adjust your grade up or down)
- Exam II was on March 30th, 2023 during lecture
 - Still a few people who haven't taken it yet...so won't discuss

▶ Labs

- Lab 5 (Process Lab) was due on March 30th, 2023 @ 11:59 PM EST

▶ Projects

- Project III: Late submission closed on March 27th, 2023 @11:59 PM EST
 - Remember to schedule check-off meetings if you haven't already
- Project IV was released on March 30th, 2023
 - **Due: April 10th, 2023 @ 11:59 PM EST**

▶ Poll Everywhere

- www.polllev.com/shinwookim908
- Solutions to recitation questions will be posted on website

PEV: Signals

When poll is active, respond at pollev.com/shinwookim908

Which of the following are TRUE about signals?

SIGKILL can be ignored

Users can use custom signals, like SIGUSR1

SIGSEGV happens when a child process terminates.

None of the above




Powered by  Poll Everywhere

Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

Project IV

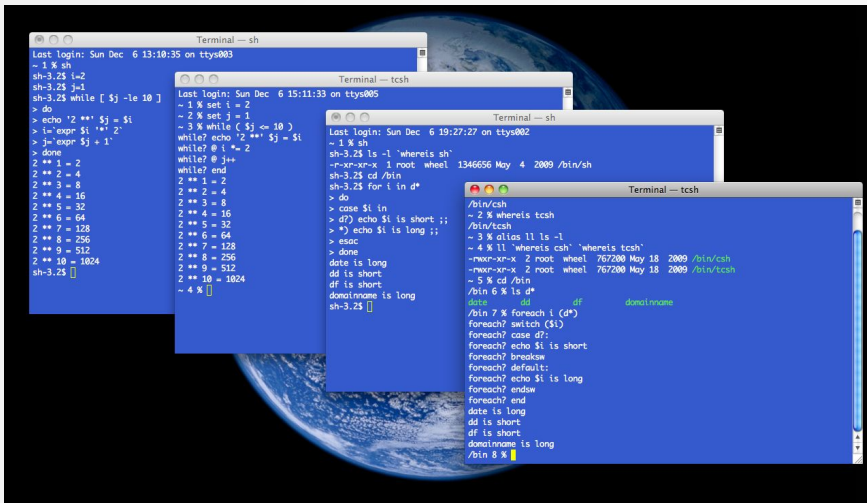
Writing your own shell



The shell

is the outermost layer of the operating system

What's a shell?

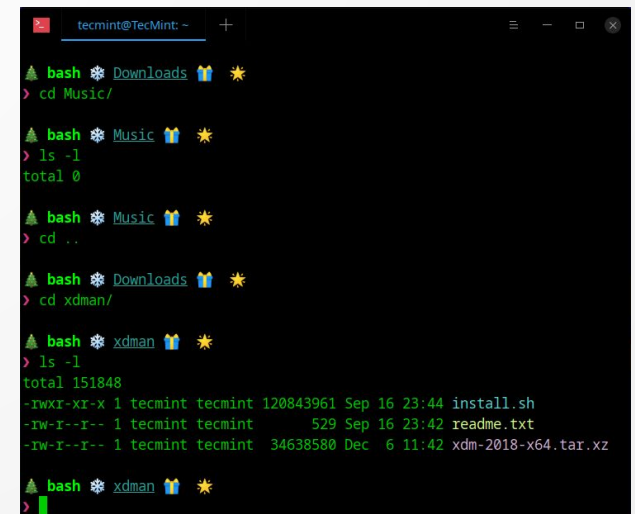
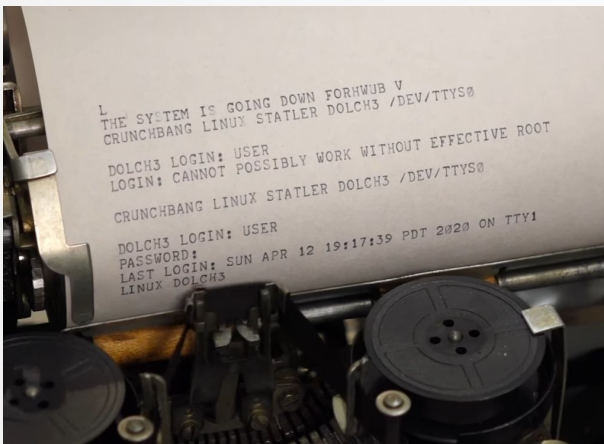


The image shows four overlapping terminal windows. The top-left window is titled 'Terminal - sh' and shows a script that prints numbers 1 through 10. The top-right window is titled 'Terminal - tcsh' and shows a script that prints numbers 1 through 10. The bottom-left window is titled 'Terminal - sh' and shows a script that prints numbers 1 through 10. The bottom-right window is titled 'Terminal - tcsh' and shows a script that prints numbers 1 through 10.

- ▶ It's the “*command line*”
- ▶ A ***shell*** is an application program that runs programs on behalf of the user.
- ▶ Typically a shell is a program that
 1. Repeatedly prints a prompt
 2. Waits for a *command line* on **stdin**
 3. Carries out some action (as directed by the contents of the command line)
- ▶ A ***Read*** → ***Evaluate*** → ***Print*** loop (***REPL***)

Some terminology

- ▶ A **shell** is a user interface for accessing an computer system
- ▶ Most often the user interacts with the shell using a **command-line interface (CLI)**.
- ▶ The **terminal** is a program that opens a graphical window and lets you interact with the shell.
 - Actually this is a **terminal emulator** or **virtual console**
 - Technically, terminals are physical machines that provides an interface with a larger machine
 - Teletypewriters
 - Video display terminals
- ▶ In reality, all these terms are *more or less* used interchangeably.



Many different shells, including your very own!

- ▶ There are various different shells that you can use.
 - `sh` – Original Unix shell (Stephen Bourne, AT&T Bell Labs, 1977)
- ▶ Most common is the Bourne-Again shell (`bash`)
 - Preinstalled with most Linux distributions
 - It's the one that's installed on Thoth
 - Just another program → `/bin/bash`
- ▶ Some others include:
 - Z-shell (`zsh`) → `/bin/zsh`
 - Comes preinstalled for modern MacOS, modern Linux distributions
 - PowerShell, COMMAND.COM
 - For Windows
 - Not a Unix-Shell
 - `fish/csh`, and much more
- ▶ For project IV, you will implement your very own shell
 - Primitive, yet still functional
 - *It accomplishes all that needs to be done*

msh specification

Hopefully you can come up with a good name for your shell that ends with “-sh”

Your shell should:

- ▶ Print a prompt: “>”
- ▶ Read user input
 - The command line input by the user consists of a *name* and zero or more arguments (delimited by spaces)

```
> ls                # command: ls; arguments: ls
> ls -a             # command: ls; arguments: ls, -a
> exit              # command: exit; arguments: exit
> load better_ls    # command: load; arguments load, better_ls
```

msh specification

Your shell should:

► Support built-in commands

- `exit`: The shell should exit upon receiving this command
- `load`: The shell should dynamically load a plugin and initialize it

► Support extensioning built-in commands via plugins

- Plugin Interface:
 - `int initialize()`
 - Returns 0 on success
 - `int run(char **argv)`
 - `argv`: array of Strings terminated by `NULL`
 - `argv = {"ls", "-a", NULL}`
 - Returns 0 on success
- Throw error message if plugin could not be loaded
`Error: Plugin <plugin> initialization failed!`
- Once loaded, user should be able to run the extended functionality by invoking the plugin's name

msh specification

Your shell should:

- ▶ Support extensioning built-in commands via plugins

```
> broken_better_ls    # Not loaded
> load broken_better_ls
Error: Plugin broken_better_ls initialization failed!
> broken_better_ls    # Still not loaded
> better_ls           # Not loaded
> load better_ls      # Success
> better_ls           # Loaded
msh      msh.c        better_ls.c    better_ls.so
>
```

msh specification

Your shell should:

- ▶ Instantiating other executables

```
shk148@thoth $ ./msh
```

```
> vim better_ls.c
```

```
> gcc better_ls.c -o better_ls.so -shared
```

```
> load better_ls
```

```
> better_ls
```

```
msh      msh.c      better_ls.c  better_ls.so
```

```
> exit
```

```
shk148@thoth $
```

msh specification

Your shell should:

- ▶ Instantiating other executables

```
shk148@thoth $ ./msh
```

```
> vim better_ls.c
```

```
> gcc better_ls.c -o better_ls.so -shared
```

```
> load better_ls
```

```
> better_ls
```

```
msh      msh.c      better_ls.c  better_ls.so
```

```
> exit
```

```
shk148@thoth $
```

msh specification limitations

To simplify your implementation, testing will be limited to:

1. Commands will have a maximum size of 200 characters
2. Program names and arguments will have a maximum size of 20 characters
3. There will be at most 20 arguments
4. Your shell need only support loading upto 10 plugins

Building the shell: ~~Skeleton~~ Shelleton

```
int main(){
    while (TRUE) When do we break out of this loop?
    { /* Infinite Loop for REPL */
        PrintCommandPrompt()
        cmdLine = readFromStdIn();
        cmd = parseCommand(cmdLine);
        If (cmd is BuiltInCommand) {executeBuiltInCommand(cmd)};
        Else
        { If the command not a built-in command, we should check if it's a name of an executable file
            fork()
            // Child process should run the executable
            What should the parent process do while the child process is running?
        }
    }
}
```

...This is just one approach to building your shell

Review: C Strings

What does the following program output?

```
#include <stdio.h>

int main ()
{
    char str[25] = "Computersystems";
    printf ("%s", str + 8);
    return 0;
}
```


PEV: C Strings

What does the following code output?

Review: C Strings

What does the following program output?

```
#include <stdio.h>
int main ()
{
    char str[25] = "Computersystems";
    printf ("%s", str + 8);
    return 0;
}
```




Building the shell: Reading and Parsing Input

Built-in command or path to another executable

Command line arguments

\$ ls -l -a /usr

- ▶ A command goes in →  → a process comes out
 - A shell, at its simplest, is a program that reads input from the user and tries to execute commands.
- ▶ We can read in a line of input using `scanf()`
- ▶ Given a user input, we need to categorize it as
 - Built in command or
 - Name of an executable
- ▶ But before we can interpret the input, we need to tokenize it

"ls -l -a /usr" /* delimited by ' */
⇒ {"ls", "-l", "-a", "/usr"}

Review: Reading Input

What does the following code do?

```
scanf("%7s",ch);
```

- a) read string with minimum 7 characters.
- b) read string with maximum 7 characters
- c) read string exactly to 7 characters
- d) read string with any number of characters

PEV: Reading Input

🌐 When poll is active, respond at pollev.com/shinwookim908

What does `scanf ("%7s", ch)` do?

Review: Reading Input

What does the following code do?

```
scanf ("%7s", ch);
```

- a) read string with minimum 7 characters.
- b) read string with maximum 7 characters
- c) read string exactly to 7 characters
- d) read string with any number of characters

Read string with minimum 7 characters.

Read string with maximum 7 characters

Read string exactly to 7 characters

Read string with any number of characters

None of the above



- ▶ The `strtok()` function can help tokenize strings
- ▶ `#include <string.h>`
- ▶ `char *strtok(char *str, const char *delim);`
 - Breaks string `str` into a series of tokens using the delimiter `delim`.
 - Returns a pointer to the next token, or NULL if there are no more tokens.
- ▶ Called in one of two ways:
 1. `strtok(str, d)` // starts processing a new string
 2. `strtok(NULL, d)` // continue processing a string

A strtok() example

```
$ ./strtok_example
```

```
I
```

```
#include <stdio.h>
#include <string.h>
int main(){
    char str[] = "I:love-programming";
    char delim[] = "-:.";
    char *token;
    token = strtok(str, delim);
    printf("%s\n", token); ← What will be printed?
    return 0;
}
```

A strtok() example

```
#include <stdio.h>
#include <string.h>

int main(){
    char str[] = "I:love-programming";
    char delim[] = "-:.";
    char *token;
    token = strtok(str, delim);
    printf("%s\n", token);
    token = strtok(str, delim);
    printf("%s\n", token); ←————— What will be printed?
    return 0;
}
```

```
$ ./strtok_example
```

```
I
```

```
I
```

🤔 But the second token should be "love"

A strtok() example

```
#include <stdio.h>
#include <string.h>

int main(){
    char str[] = "I:love-programming";
    char delim[] = "-:.";
    char *token;
    token = strtok(str, delim);
    printf("%s\n", token);
    token = strtok(NULL, delim);
    printf("%s\n", token);
    return 0;
}
```

```
$ ./strtok_example
```

```
I
love
```

How can we print the remaining tokens?

← What will be printed?

A strtok() example

```
char* s = "See the red fox";
```

char* s =

S	e	e		t	h	e		r	e	d		f	o	x	\0	...
---	---	---	--	---	---	---	--	---	---	---	--	---	---	---	----	-----

```
char* t = strtok(s, " ");
```

char* s =

S	e	e	\0	t	h	e		r	e	d		f	o	x	\0	...
---	---	---	----	---	---	---	--	---	---	---	--	---	---	---	----	-----

t

```
char* t = strtok(NULL, " ");
```

char* s =

S	e	e	\0	t	h	e	\0	r	e	d		f	o	x	\0	...
---	---	---	----	---	---	---	----	---	---	---	--	---	---	---	----	-----

t

```
char* t = strtok(NULL, " ");
```

char* s =

S	e	e	\0	t	h	e	\0	r	e	d		f	o	x	\0	...
---	---	---	----	---	---	---	----	---	---	---	--	---	---	---	----	-----

```
char* t = strtok(NULL, " ");
```

t

char* s =

S	e	e	\0	t	h	e	\0	r	e	d	\0	f	o	x	\0	...
---	---	---	----	---	---	---	----	---	---	---	----	---	---	---	----	-----

t

```
char* t = strtok(NULL, " ");    t → NULL
```

► **strtok()** changes the string that has been parsed!

idem·po·tent

- ▶ The `strtok()` function exhibits some weird behavior
 - `strtok()` changes the string that has been parsed
 - Replacing the character in place with a null terminator (`'\0'`)
- ▶ `strtok()` produces different results when called multiple times
 - It's a **non-idempotent** function
 - Which has **side effects**.
- ▶ In comparison, functions that have no side effects are called **idempotent**.

```
x = 2; // Assignment operations are
x = 2; // idempotent
x = 2;
x = 2; // Calling it multiple times
x = 2; // always produces the same result
```

man strtok #notes-and-bugs

- ▶ Be cautious when using these functions. If you do use them, note that:
 - These functions modify their first argument.
 - These functions cannot be used on constant strings.
 - The identity of the delimiting byte is lost.
- ▶ For instance, if you try
 - `strtok("String Constant", delim)`
 - Segmentation fault! (attempting to write to a literal)

Still unsure? Read the man pages!

`$ man strtok`

- ▶ What arguments does the function take?
 - read **SYNOPSIS**
- ▶ What does the function do?
 - read **DESCRIPTION**
- ▶ What does the function return?
 - read **RETURN VALUES**
- ▶ What errors can the function fail with?
 - read **ERRORS**
- ▶ Is there anything I should watch out for?
 - read **NOTES**
- ▶ I want an example
 - read **EXAMPLES**

strtok() vs strsep()

- ▶ Alternatively, you can use `strsep()`
- ▶ A *replacement* for `strtok()`
- ▶ But not all C versions support it
 - For instance, ANSI-C does not support `strtok()`
 - Hence, it is *less portable*
- ▶ You may use either `strsep()` or `strtok()` in this project
 - Read the documentation (man pages) to see how each work!

Building the shell: Executing command

- ▶ Once we've tokenized the input, we can use standard C-string functions to *compare*
 - `strcmp()` and friends
- ▶ If the keyword matches a built-in command
 - Run it!
 - Some functionalities may require dynamically loading *plugins*
 - Just as you did for lab 5
- ▶ If the keyword is unknown,
 - It's probably the name of an executable
 - So run it!
 - `fork()` and friends

Building the shell: Executing command

- ▶ Once we've tokenized the input, we can use standard C-string functions to *compare*
 - `strcmp()` and friends
- ▶ If the keyword matches a built-in command
 - Run it!
 - Refer to lab 5 on how to dynamically load plugins
- ▶ If the keyword is unknown,
 - It's probably the name of an executable
 - So run it!
 - `fork()` and friends

Building the shell: Executing command

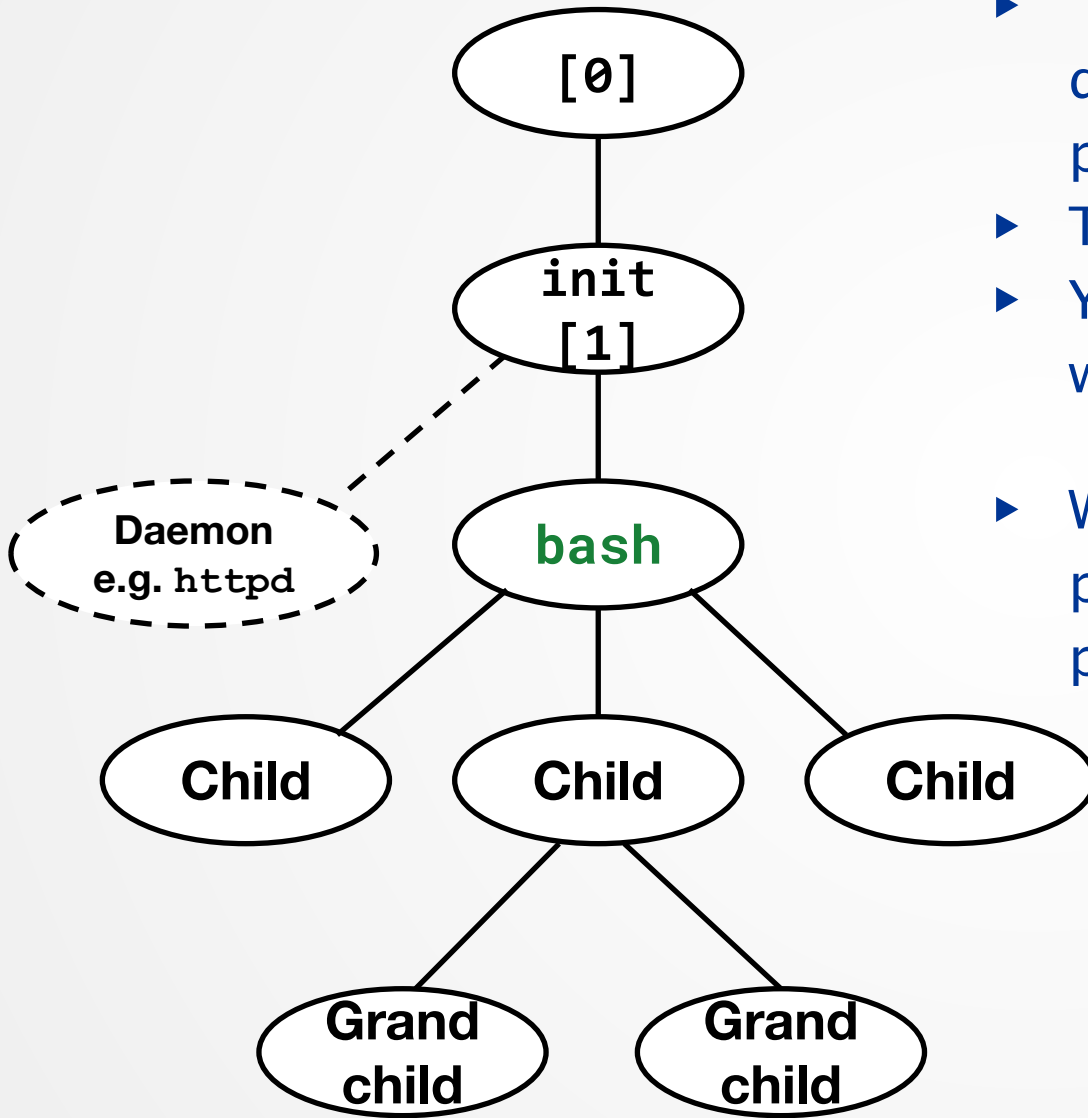

- ▶ Once we've tokenized the input, we can use standard C-string functions to *compare*
 - `strcmp()` and friends
- ▶ If the keyword matches a built-in command
 1. `exit` ⇒ Exit the program
 2. `load` ⇒ Dynamically load plugins (just like lab 5)

■

Building the shell: Executing command

- ▶ Once we've tokenized the input, we can use standard C-string functions to *compare*
- ▶ If the keyword does not match a built-in command
 - Check if it's a plugin...and run it!
 - If it's not a plugin...must be an executable name!

UNIX Process Hierarchy



- ▶ UNIX/Linux has an interesting design: every application is a child process.
- ▶ The root is the init task.
- ▶ Your shell spawns child processes when you ask to run a command.
 - Using `fork()` and `exec()`!
- ▶ When your own application spawns a process, it spawns its own child process using `fork()`
 - If your app exits before the child, the child becomes an orphan 😞
 - Orphan processes are *adopted* by the root

Building the shell: Executing command

- ▶ If the keyword does not match a built-in command
 - and If it's not a plugin...must be an executable name!
- ▶ `fork()`, `exec*`, and their friends!
 - Make sure to use the correct `exec*` function
 - Make sure to correctly pass in arguments

Implementation Challenges & Hints

1. Since our shell needs to support dynamically loading multiple plugins
 - Devise some data structure to store them
 - Create helper functions to add and access plugins
 2. When multiprocessing with `fork()`s
 - Think about the order in which processes need to run
 - Does a process need to wait for another?
 3. String parsing is weird and hard
 - Especially since the standard functions exhibits odd behavior
 - Carefully read the documentation
 - Verify output before moving onto next step
- Since this project requires access to many standard library functions, we highly recommend developing on Thoth or another Linux machine
- And plan for outages!
 - i. Back-up frequently (to your local machine)

Works Referred

- ▶ Creative Commons photography courtesy of Arnold Reinhold and technikum29 via the Wikimedia Foundations
- ▶ `strtok()` examples adapted from Weber State University