



# Project 4:

# Writing your own shell

---

**Shinwoo Kim**

Teaching Assistant

shinwookim@pitt.edu

<https://www.pitt.edu/~shk148/>

Spring 2023, Term 2234

Friday 12 PM Recitation

5502 Sennott Square

Mar 2nd, 2023

# Course News!

---

## ▶ Exams

- Exam I grades were returned on March 24th, 2023
  - Check your email for class statistics
  - Request regrades if needed (this may adjust your grade up or down)
- Exam II was on March 30th, 2023 during lecture
  - Still a few people who haven't taken it yet...so won't discuss

## ▶ Labs

- Lab 5 (Process Lab) was due on March 30th, 2023 @ 11:59 PM EST

## ▶ Projects

- Project III: Late submission closed on March 27th, 2023 @11:59 PM EST
  - Remember to schedule check-off meetings if you haven't already
- Project IV was released on March 30th, 2023
  - **Due: April 10th, 2023 @ 11:59 PM EST**

## ▶ Poll Everywhere

- [www.polllev.com/shinwookim908](http://www.polllev.com/shinwookim908)
- Solutions to recitation questions will be posted on website

# PEV: Signals

---

🌐 When poll is active, respond at [pollev.com/shinwookim908](https://pollev.com/shinwookim908)

## Which of the following are TRUE about signals?

SIGKILL can be ignored

Users can use custom signals, like SIGUSR1

SIGSEGV happens when a child process terminates.

None of the above




Powered by  Poll Everywhere

Start the presentation to see live content. For screen share software, share the entire screen. Get help at [pollev.com/app](https://pollev.com/app)

# Project IV

---

Writing your own shell

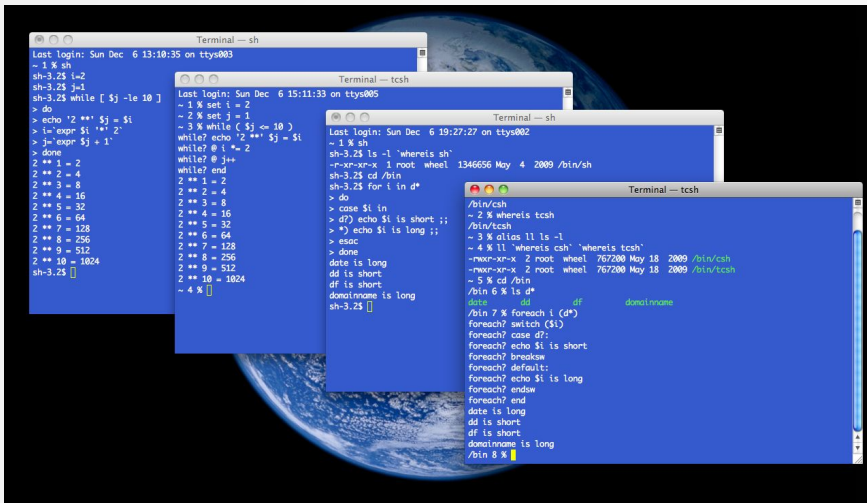


# The shell

---

is the outermost layer of the operating system

# What's a shell?

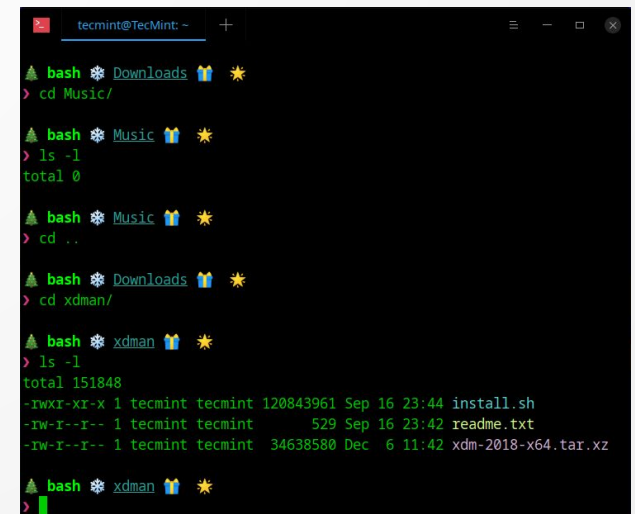
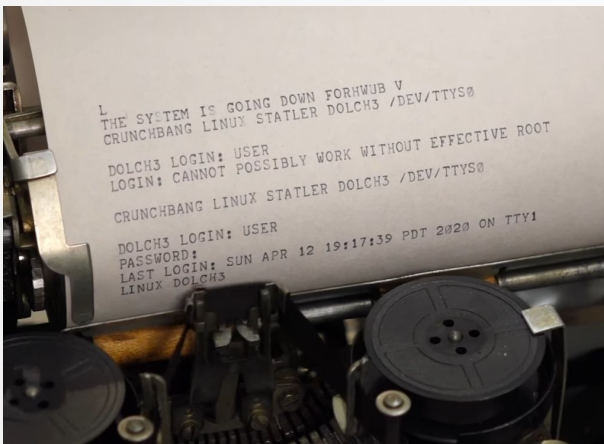


The image shows four overlapping terminal windows. The top-left window is titled 'Terminal - sh' and shows a script that prints numbers 1 through 10. The top-right window is titled 'Terminal - tcsh' and shows a script that prints numbers 1 through 10. The bottom-left window is titled 'Terminal - sh' and shows a script that prints numbers 1 through 10. The bottom-right window is titled 'Terminal - tcsh' and shows a script that prints numbers 1 through 10.

- ▶ It's the “*command line*”
- ▶ A ***shell*** is an application program that runs programs on behalf of the user.
- ▶ Typically a shell is a program that
  1. Repeatedly prints a prompt
  2. Waits for a *command line* on **stdin**
  3. Carries out some action (as directed by the contents of the command line)
- ▶ A ***Read*** → ***Evaluate*** → ***Print*** loop (***REPL***)

# Some terminology

- ▶ A **shell** is a user interface for accessing an computer system
- ▶ Most often the user interacts with the shell using a **command-line interface (CLI)**.
- ▶ The **terminal** is a program that opens a graphical window and lets you interact with the shell.
  - Actually this is a **terminal emulator** or **virtual console**
  - Technically, terminals are physical machines that provides an interface with a larger machine
    - Teletypewriters
    - Video display terminals
- ▶ In reality, all these terms are *more or less* used interchangeably.



# Many different shells, including your very own!

---

- ▶ There are various different shells that you can use.
  - `sh` – Original Unix shell (Stephen Bourne, AT&T Bell Labs, 1977)
- ▶ Most common is the Bourne-Again shell (`bash`)
  - Preinstalled with most Linux distributions
    - It's the one that's installed on Thoth
  - Just another program → `/bin/bash`
- ▶ Some others include:
  - Z-shell (`zsh`) → `/bin/zsh`
    - Comes preinstalled for modern MacOS, modern Linux distributions
  - PowerShell, COMMAND.COM
    - For Windows
    - Not a Unix-Shell
  - `fish/csh`, and much more
- ▶ For project IV, you will implement your very own shell
  - Primitive, yet still functional
  - *It accomplishes all that needs to be done*



# msh specification

*Hopefully you can come up with a good name for your shell that ends with “-sh”*

---

## Your shell should:

- ▶ Print a prompt: “>”
- ▶ Read user input
  - The command line input by the user consists of a *name* and zero or more arguments (delimited by spaces)

```
> ls                # command: ls; arguments: ls
> ls -a             # command: ls; arguments: ls, -a
> exit              # command: exit; arguments: exit
> load better_ls    # command: load; arguments load, better_ls
```

# msh specification

---

## Your shell should:

### ► Support built-in commands

- `exit`: The shell should exit upon receiving this command
- `load`: The shell should dynamically load a plugin and initialize it

### ► Support extensioning built-in commands via plugins

- Plugin Interface:
  - `int initialize()`
    - Returns 0 on success
  - `int run(char **argv)`
    - `argv`: array of Strings terminated by `NULL`
      - `argv = {"ls", "-a", NULL}`
    - Returns 0 on success
- Throw error message if plugin could not be loaded  
`Error: Plugin <plugin> initialization failed!`
- Once loaded, user should be able to run the extended functionality by invoking the plugin's name

# msh specification

---

Your shell should:

- ▶ Support extensioning built-in commands via plugins

```
> broken_better_ls    # Not loaded
> load broken_better_ls
Error: Plugin broken_better_ls initialization failed!
> broken_better_ls    # Still not loaded
> better_ls           # Not loaded
> load better_ls      # Success
> better_ls           # Loaded
msh      msh.c        better_ls.c    better_ls.so
>
```

# msh specification

---

Your shell should:

- ▶ Allow for instantiating other executables *and pass in arguments*

```
shk148@thoth $ ./msh
```

```
> vim better_ls.c
```

```
> gcc better_ls.c -o better_ls.so -shared
```

```
> load better_ls
```

```
> better_ls
```

```
msh      msh.c      better_ls.c  better_ls.so
```

```
> exit
```

```
shk148@thoth $
```

# msh specification limitations

---

To simplify your implementation, testing will be limited to:

1. Commands will have a maximum size of 200 characters
2. Program names and arguments will have a maximum size of 20 characters
3. There will be at most 20 arguments
4. Your shell need only support loading upto 10 plugins

# Building the shell: ~~Skeleton~~ Shelleton

---

```
int main(){
    while (TRUE) When do we break out of this loop?
    { /* Infinite Loop for REPL */
        PrintCommandPrompt()
        cmdLine = readFromStdIn();
        cmd = parseCommand(cmdLine);
        If (cmd is BuiltInCommand) {executeBuiltInCommand(cmd)};
        Else
        { If the command not a built-in command, we should check if it's a name of an executable file
            fork()
            // Child process should run the executable
            What should the parent process do while the child process is running?
        }
    }
}
```

*...This is just one approach to building your shell*

# Review: C Strings

---

What does the following program output?

```
#include <stdio.h>

int main ()
{
    char str[25] = "Computersystems";
    printf ("%s", str + 8);
    return 0;
}
```

# PEV: C Strings

---

## What does the following code output?

### Review: C Strings

---

What does the following program output?

```
#include <stdio.h>
int main ()
{
    char str[25] = "Computersystems";
    printf ("%s", str + 8);
    return 0;
}
```






# Building the shell: Reading and Parsing Input

*Built-in command or path to another executable*

*Command line arguments*

\$ ls -l -a /usr

- ▶ A command goes in →  → a process comes out
  - A shell, at its simplest, is a program that reads input from the user and tries to execute commands.
- ▶ We can read in a line of input using `fgets()`
- ▶ Given a user input, we need to categorize it as
  - Built in command or
  - Name of an executable
- ▶ But before we can interpret the input, we need to tokenize it

"ls -l -a /usr" /\* delimited by ' \*/  
⇒ {"ls", "-l", "-a", "/usr"}

- ▶ The `strtok()` function can help tokenize strings
- ▶ `#include <string.h>`
- ▶ `char *strtok(char *str, const char *delim);`
  - Breaks string `str` into a series of tokens using the delimiter `delim`.
  - Returns a pointer to the next token, or NULL if there are no more tokens.
- ▶ Called in one of two ways:
  1. `strtok(str, d)` // starts processing a new string
  2. `strtok(NULL, d)` // continue processing a string

# A strtok() example

---

```
$ ./strtok_example
```

```
I
```

```
#include <stdio.h>
#include <string.h>
int main(){
    char str[] = "I:love-programming";
    char delim[] = "-:.";
    char *token;
    token = strtok(str, delim);
    printf("%s\n", token); ← What will be printed?
    return 0;
}
```

# A strtok() example

```
#include <stdio.h>
#include <string.h>

int main(){
    char str[] = "I:love-programming";
    char delim[] = "-:.";
    char *token;
    token = strtok(str, delim);
    printf("%s\n", token);
    token = strtok(str, delim);
    printf("%s\n", token); ←————— What will be printed?
    return 0;
}
```

```
$ ./strtok_example
```

```
I
I
```

🤔 But the second token should be "love"

# A strtok() example

```
#include <stdio.h>
#include <string.h>

int main(){
    char str[] = "I:love-programming";
    char delim[] = "-:.";
    char *token;
    token = strtok(str, delim);
    printf("%s\n", token);
    token = strtok(NULL, delim);
    printf("%s\n", token);
    return 0;
}
```

```
$ ./strtok_example
```

```
I
love
```

*How can we print the remaining tokens?*

← What will be printed?

# A strtok() example

```
char* s = "See the red fox";
```

char\* s =

S	e	e		t	h	e		r	e	d		f	o	x	\0	...
---	---	---	--	---	---	---	--	---	---	---	--	---	---	---	----	-----

```
char* t = strtok(s, " ");
```

char\* s =

S	e	e	\0	t	h	e		r	e	d		f	o	x	\0	...
---	---	---	----	---	---	---	--	---	---	---	--	---	---	---	----	-----

t

```
char* t = strtok(NULL, " ");
```

char\* s =

S	e	e	\0	t	h	e	\0	r	e	d		f	o	x	\0	...
---	---	---	----	---	---	---	----	---	---	---	--	---	---	---	----	-----

t

```
char* t = strtok(NULL, " ");
```

char\* s =

S	e	e	\0	t	h	e	\0	r	e	d	\0	f	o	x	\0	...
---	---	---	----	---	---	---	----	---	---	---	----	---	---	---	----	-----

```
char* t = strtok(NULL, " ");
```

t

char\* s =

S	e	e	\0	t	h	e	\0	r	e	d	\0	f	o	x	\0	...
---	---	---	----	---	---	---	----	---	---	---	----	---	---	---	----	-----

```
char* t = strtok(NULL, " ");
```

t

t → NULL

► strtok() changes the string that has been parsed!

# idem·po·tent

---

- ▶ The `strtok()` function exhibits some weird behavior
  - `strtok()` changes the string that has been parsed
  - Replacing the character in place with a null terminator ( `'\0'` )
- ▶ `strtok()` produces different results when called multiple times
  - It's a **non-idempotent** function
    - Which has **side effects**.
- ▶ In comparison, functions that have no side effects are called **idempotent**.

```
x = 2; // Assignment operations are
x = 2; // idempotent
x = 2;
x = 2; // Calling it multiple times
x = 2; // always produces the same result
```

- ▶ Be cautious when using these functions. If you do use them, note that:
  - These functions modify their first argument.
  - These functions cannot be used on constant strings.
  - The identity of the delimiting byte is lost.
- ▶ For instance, if you try
  - `strtok("String Constant", delim)`
  - Segmentation fault! (attempting to write to a literal)



# Still unsure? Read the man pages!

---

`$ man strtok`

- ▶ What arguments does the function take?
  - read **SYNOPSIS**
- ▶ What does the function do?
  - read **DESCRIPTION**
- ▶ What does the function return?
  - read **RETURN VALUES**
- ▶ What errors can the function fail with?
  - read **ERRORS**
- ▶ Is there anything I should watch out for?
  - read **NOTES**
- ▶ I want an example
  - read **EXAMPLES**
  - <https://pitt.edu/~shk148/teaching/CS0449-2234/code/strtok.c.html>

# strtok() vs strsep()

---

- ▶ Alternatively, you can use `strsep()`
- ▶ A *replacement* for `strtok()`
- ▶ But not all C versions support it
  - For instance, ANSI-C does not support `strtok()`
  - Hence, it is *less portable*
- ▶ You may use either `strsep()` or `strtok()` in this project
  - Read the documentation (man pages) to see how each work!

# Building the shell: Executing command

---

- ▶ Once we've tokenized the input, we can use standard C-string functions to *compare*
  - `strcmp()` and friends
- ▶ If the keyword matches a built-in command
  - Run it!
  - Some functionalities may require dynamically loading *plugins*
    - Just as you did for lab 5
- ▶ If the keyword is unknown,
  - It's probably the name of an executable
  - So run it!
    - `fork()` and friends
      - `exec*()`
        - `wait()`

# Building the shell: Executing command

---

- ▶ Once we've tokenized the input, we can use standard C-string functions to *compare*
  - `strcmp()` and friends
- ▶ If the keyword matches a built-in command
  - Run it!
  - Refer to lab 5 on how to dynamically load plugins
- ▶ If the keyword is unknown,
  - It's probably the name of an executable
  - So run it!
    - `fork()` and friends

# Building the shell: Executing command

---

```
{"ls", "-l", "-a", "/usr"}
```

- ▶ Once we've tokenized the input, we can use standard C-string functions to *compare*
  - `strcmp()` and friends
- ▶ If the keyword matches a built-in command
  1. `exit` ⇒ Exit the program
  2. `load` ⇒ Dynamically load plugins (just like lab 5)
    - Since our shell needs to support dynamically loading multiple plugins
      - Devise some data structure to store them
      - Create helper functions to add and access plugins

# Building the shell: Executing command

---

- ▶ If the keyword does not match a built-in command
- ▶ Check if it's a plugin
  - and run it
- ▶ If it's not a plugin
  - It must be an executable name
  - `fork()`, `exec*()`, and their friends!
    - Make sure to use the correct `exec*()` function
    - And correctly pass in arguments

# Implementation Hints

---

## 1. When multiprocessing with `fork()`s

- Think about the order in which processes need to run
- Does a process need to wait for another?

## 2. String parsing is weird and hard

- Especially since the standard functions exhibits odd behavior
- Carefully read the documentation
- Verify output before moving onto next step

## 3. There is a lot to program

- Break your program down into smaller functions
- `readInput()`, `parseInput()`, `runBuiltIn()`, ...
- To pass values between functions, you have to store them in the heap!

## ➤ Since this project requires access to many standard library functions, we highly recommend developing on Thoth or another Linux machine

- And plan for outages!
  - Back-up frequently (to your local machine)

# Implementation Challenges

---

## 1. This project ties in everything you've learned so far

- C programming & debugging
  - See *Lab0 (Hello lab)*
- C-Strings and standard library functions
  - See *Project I (BMP Steganography)* for a guide
- Maintaining data structures in C
  - *Lab3 (Queue lab)*
- Pointers and management of memory
  - See *Lab2 (Pointer lab)*, *Project II (Malloc)*
- Process management and dynamic loading
  - See *Lab5 (Loading and Forking)*

## 2. One common issue: *Memory leaks*

- Not maintaining pointers
- `malloc()` without `free()`
- Test your code for memory leaks using **valgrind**!



# Implementation Challenges

## 2. One common issue: *Memory leaks*

- Not maintaining pointers
- `malloc()` without `free()`
- Test your code for memory leaks using **valgrind**!

```
$ valgrind --leak-check=full --show-leak-kinds=all ./msh
```


### HEAP SUMMARY:

```
==630754==      in use at exit: 3,683 bytes in 6 blocks
==630754==    total heap usage: 8 allocs, 2 frees, 5,731 bytes allocated
==630754==
==630754== 820 (808 direct, 12 indirect) bytes in 1 blocks are definitely lost in loss record 4 of 5
==630754==    at 0x484DA83: calloc (in /usr/libexec/valgrind/vgpreload_memcheck-amd64-linux.so)
==630754==    by 0x10981F: get_user_input (luis.c:134)
==630754==    by 0x1097D7: main (luis.c:124)
==630754==
==630754== 2,050 bytes in 2 blocks are definitely lost in loss record 5 of 5
==630754==    at 0x484DA83: calloc (in /usr/libexec/valgrind/vgpreload_memcheck-amd64-linux.so)
==630754==    by 0x10983A: get_user_input (luis.c:137)
==630754==    by 0x1097D7: main (luis.c:124)
==630754==
==630754== LEAK SUMMARY:
==630754==    definitely lost: 2,858 bytes in 3 blocks
==630754==    indirectly lost: 12 bytes in 1 blocks
==630754==    possibly lost: 0 bytes in 0 blocks
==630754==    still reachable: 813 bytes in 2 blocks
==630754==    suppressed: 0 bytes in 0 blocks
```

# Debugging

---

- ▶ Debugging this project is hard
  - So many functionalities to look out for
    - So many places to go wrong
    - So many places to shoot yourself in the foot
  - Measure twice, cut once!
- ▶ This project is fairly open-ended in its implementation
  - *You should be able to explain your own code!*
  - ~~*"I wrote it and it sort of works, but I don't know why"*~~



GDB: GNU Debugger

*"the archer fish is known to shoot down bugs from low hanging plants by spitting water at them" -- Jamie Guinan*

[www.gnu.org/software/gdb/](http://www.gnu.org/software/gdb/)



Valgrind Memcheck

*"hunting down heap memory errors with...origami?"*

[valgrind.org](http://valgrind.org)

# Works Referred

---

- ▶ Creative Commons photography courtesy of Arnold Reinhold and technikum29 via the Wikimedia Foundations
- ▶ `strtok()` examples adapted from Weber State University